

The background of the slide is a light gray gradient. It is decorated with numerous realistic water droplets of various sizes. Some droplets are large and prominent, while others are small and subtle. They are scattered across the slide, with a higher concentration in the top-left and bottom-right corners. Each droplet has a highlight and a shadow, giving it a three-dimensional appearance.

SCHEDULING

SCHEDULING

- When more than one process is ready to run, but **only one CPU** is available, a choice is to make
- **Part** of **OS** that does it is **scheduler**
- The algorithm it uses is **scheduling algorithm**

SCHEDULING

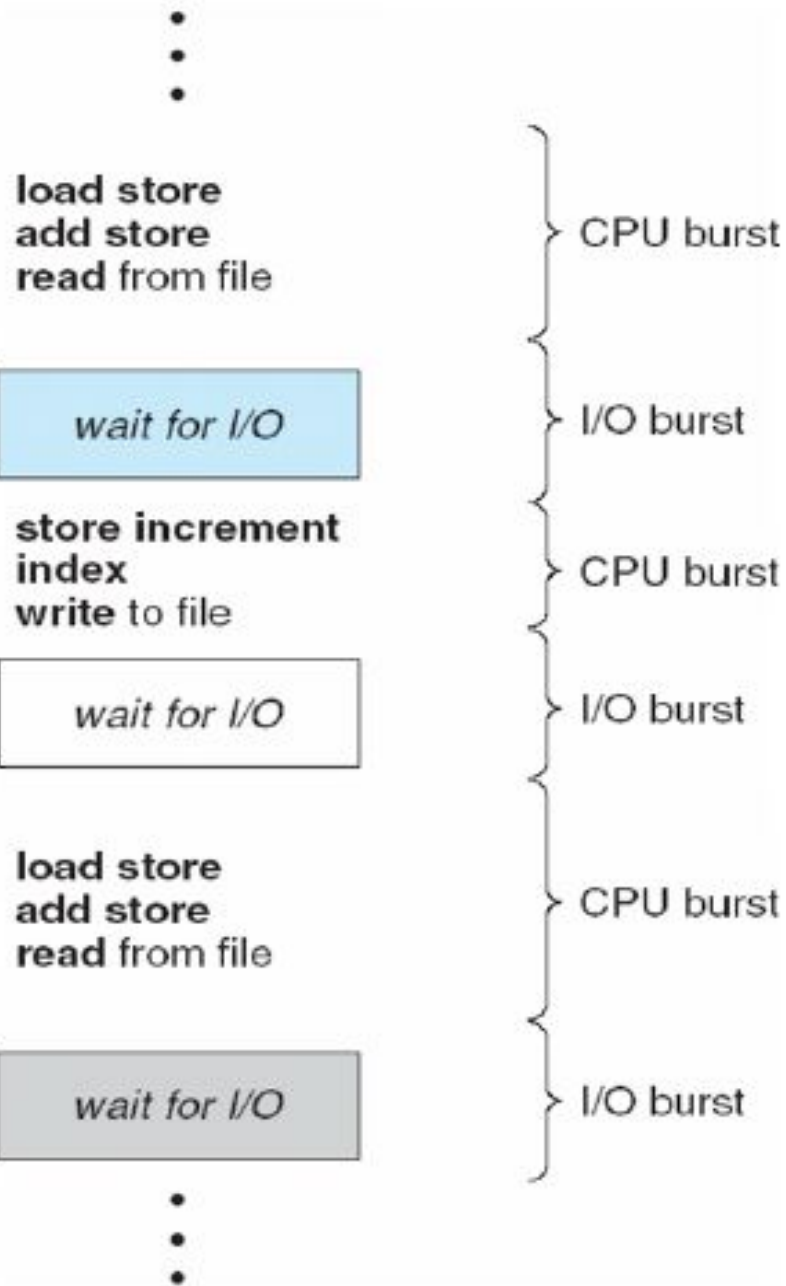
- Efficiency is needed as process switching is **costly**:
 - Switch from user mode to kernel mode
 - State of current process need to be saved
 - Memory map may be saved
 - A process is selected
 - MMU to be reloaded with memory map of new process
 - New process is started

IMPORTANCE OF SCHEDULING

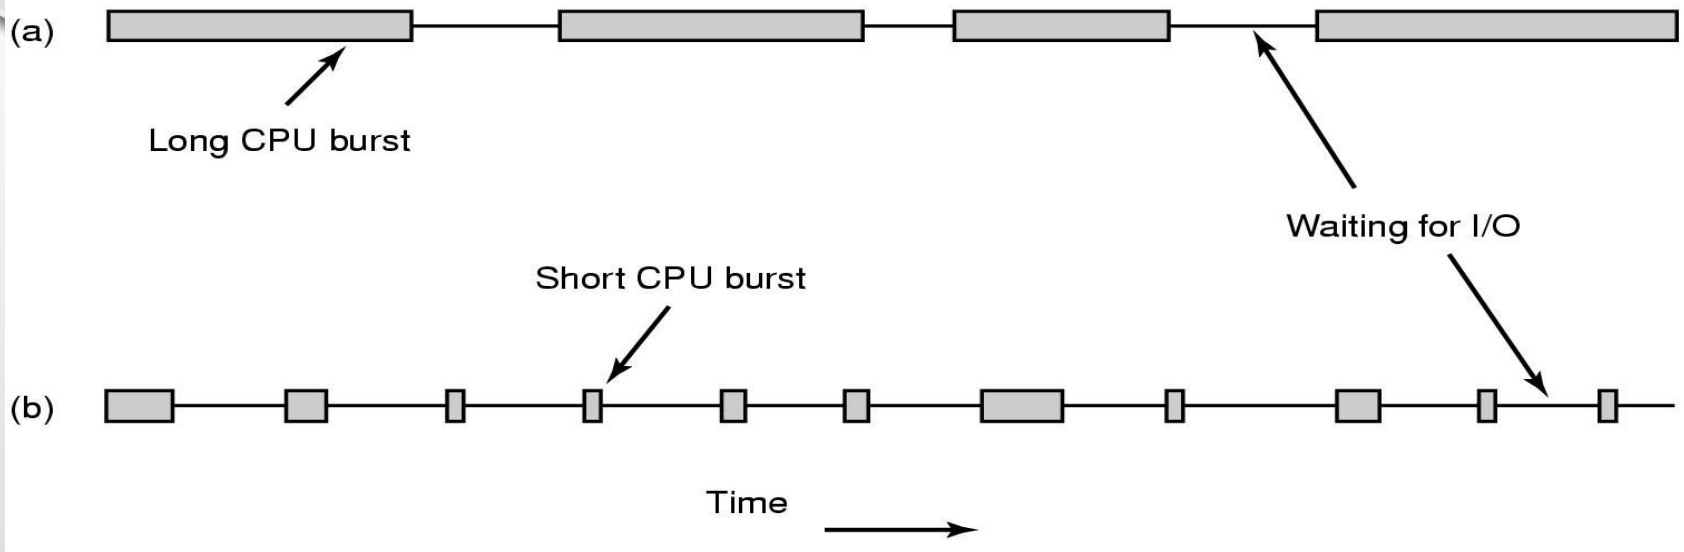
- **Good** scheduling algorithms can make a **big** difference
 - Resource utilization
 - Perceived performance & User satisfaction
 - Meeting other system goals (e.g., important tasks being taken care of immediately)

PROCESS BEHAVIOR

- Processes usually alternate **bursts** of *computing* with *I/O requests*.
- *CPU burst: the amount of time the process uses the processor before it is no longer ready*
- I/O in this sense is when a **process** enters the **blocked** state **waiting** for an external device to complete its work



PROCESS: COMPUTE AND I/O-BOUND



- a CPU-bound process (data encryption/decryption, multimedia encoding)
 - Spend **most** of the time **computing**
 - **Long** CPU bursts => infrequent I/O waits
- an I/O bound process (shell waiting for user commands)
 - Spend most of the time waiting for I/O
 - **Short** CPU bursts => frequent I/O waits
- Key factor is the **length of CPU burst** not the length of the I/O burst

PROCESS: COMPUTE AND I/O-BOUND

- As the CPUs get faster, processes tend to get more I/O bound: **WHY?**
- If a I/O bound process is ready, it should get a chance quickly.
 - Increase resource utilization

WHEN TO SCHEDULE

- When a new process is created:
 - Parent or child? Both are Ready
 - which one to run?
- When a process exits:
 - One of the ready processes should be run
- When a process blocks: Another process has to be selected to run
 - Blocking may occur for:
 - I/O
 - Semaphore

WHEN TO SCHEDULE

- When an I/O interrupt occurs:
 - In case of an interrupt of an I/O device having **completed** its work, some blocked process may now be ready
- If a h/w clock provides **periodic** interrupt: A scheduling decision can be made at each (or kth) clock interrupt

PREEMPTIVE & NON-PREEMPTIVE

Classification of **Scheduling Algorithm** depending on dealing with clock interrupt

- **Non-preemptive**: Picks a process to run and lets it run until it **blocks** or voluntarily releases the CPU. **In effect at each clock interrupt, no scheduling is done.**
- **Preemptive**: Picks a process and lets it run for a maximum of some fixed time. If still running, it is **suspended** and another is picked.
- Preemptive scheduling requires having a **clock interrupt** occur at the end of the time interval to give **control** of the CPU back to the **scheduler**

DIFFERENT SYSTEMS, DIFFERENT FOCUSES

All systems

Fairness - giving each process a fair share of the CPU

Policy enforcement - seeing that stated policy is carried out

Balance - keeping all parts of the system busy

Batch systems

Throughput - maximize jobs per hour

Turnaround time - minimize time between submission and termination

CPU utilization - keep the CPU busy all the time

Interactive systems



Response time - respond to requests quickly

Proportionality - meet users' expectations

BATCH SYSTEMS

- Users **submit** their job to the batch system
- Batch system starts user job when appropriate
- User gets notification that job is **done**
 - No interaction **in between**
- No users impatiently waiting at terminals for a **quick** response to a **short** request
- Used in business world such as Profit calculation at banks, claims processing at insurance companies...

BATCH SYSTEMS

- Common performance metrics
 - Throughput: number of jobs **completed** per hour 
 - Turnaround time: average time between the **submission** and **completion** of a job 
- Maximizing Throughput may not necessarily minimize Turnaround time (**??**)

BATCH SYSTEMS

Algorithms used:

- Non-preemptive
- Preemptive algorithms with long time periods are often acceptable
 - Reduces process switches and improves performance

Representative algorithms:

2. First Come First Serve (FCFS)
3. Shortest Job First
4. Shortest Remaining Time First

FIRST COME FIRST SERVE (FCFS)

- Process that requests the CPU FIRST is allocated the CPU FIRST.
- Also called FIFO
- **non**-preemptive
- Used in Batch Systems
- Real life analogy?
 - Transaction at Sonali Bank
- Implementation
 - FIFO queues
 - A **new** process enters the **tail** of the queue
 - The **schedule** selects from the **head** of the queue.

FCFS EXAMPLE

Process	Duration	Order	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

The final schedule:

A horizontal timeline starting at 0. A green bar labeled "P1 (24)" extends from 0 to 24. An orange bar labeled "P2 (3)" starts at 24 and ends at 27. A blue bar labeled "P3 (4)" starts at 27 and extends to the right.

P1 turnaround: 24
P2 turnaround: 27
P3 turnaround: 31

The average turnaround:
 $(24+27+31)/3 = 27.33$

Process	Duration	Order	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

The final schedule:




P1 turnaround: 24
P2 turnaround: 27
P3 turnaround: 31

The average turnaround:
 $(24+27+31)/3 = 27.33$

FCFS EXAMPLE 2

Process	Duration	Order	Arrival Time
P1	24	3	0
P2	3	1	0
P3	4	2	0

The final schedule:



P1 turnaround: 31
P2 turnaround: 3
P3 turnaround: 7

The average turnaround:
 $(31+3+7)/3 = 13.67$

ADVANTAGE

- Easy to understand and implement
- Fair for equivalent processes


PROBLEMS WITH FCFS

- Non-preemptive
- Non optimal turnaround
- **Cannot utilize** resources in **parallel**:
 - Assume **1** process CPU bounded and **many** I/O bounded processes
 - result: **Convoy effect**,
 - **low** CPU **and** I/O Device utilization
 - Why?

CONVOY EFFECT

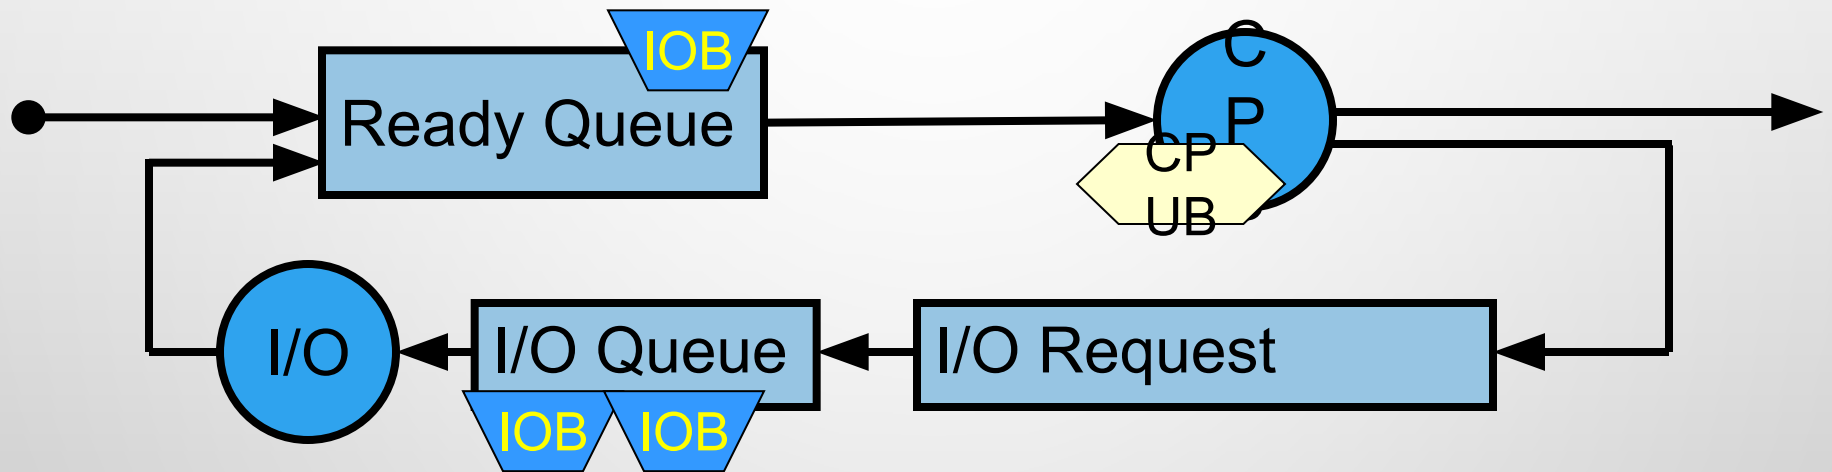
- When the CBP **uses** the CPU
 - IBPs **finish** their I/O and move into the ready queue, **waiting** for the CPU
 - the **I/O** devices are **idle**
- When the CBP finally relinquishes the CPU,
 - CBP moves to an I/O device
 - the IBPs pass through the CPU **quickly** and move back to the I/O queues
 - the CPU is **idle**
- The cycle **repeats** itself when the CBP gets back to the ready queue

CONVOY EFFECT



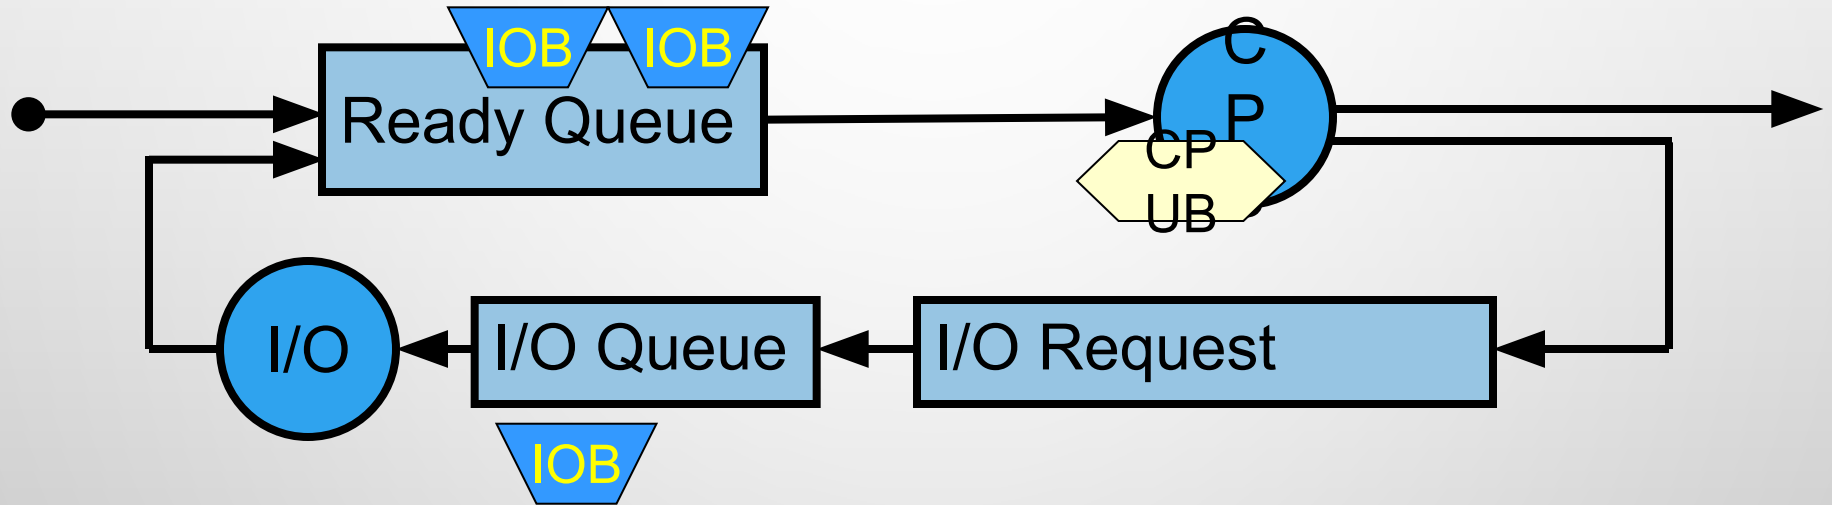
CPU is running CPUB

CONVOY EFFECT



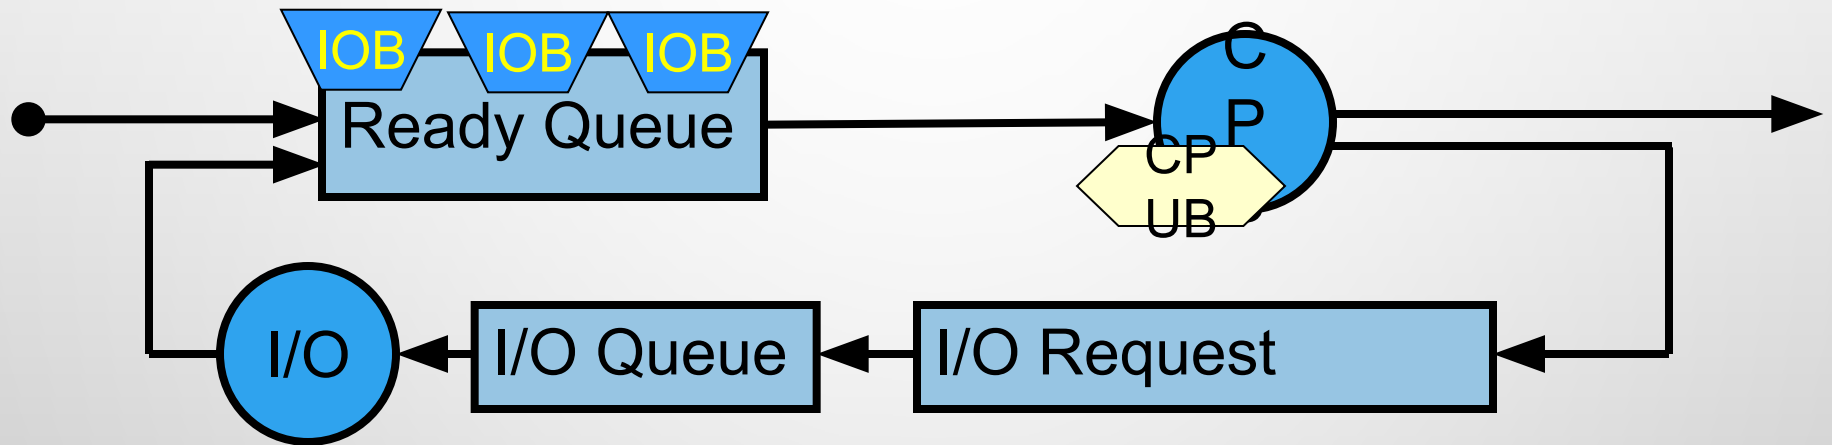
CPU is running CPUB

CONVOY EFFECT



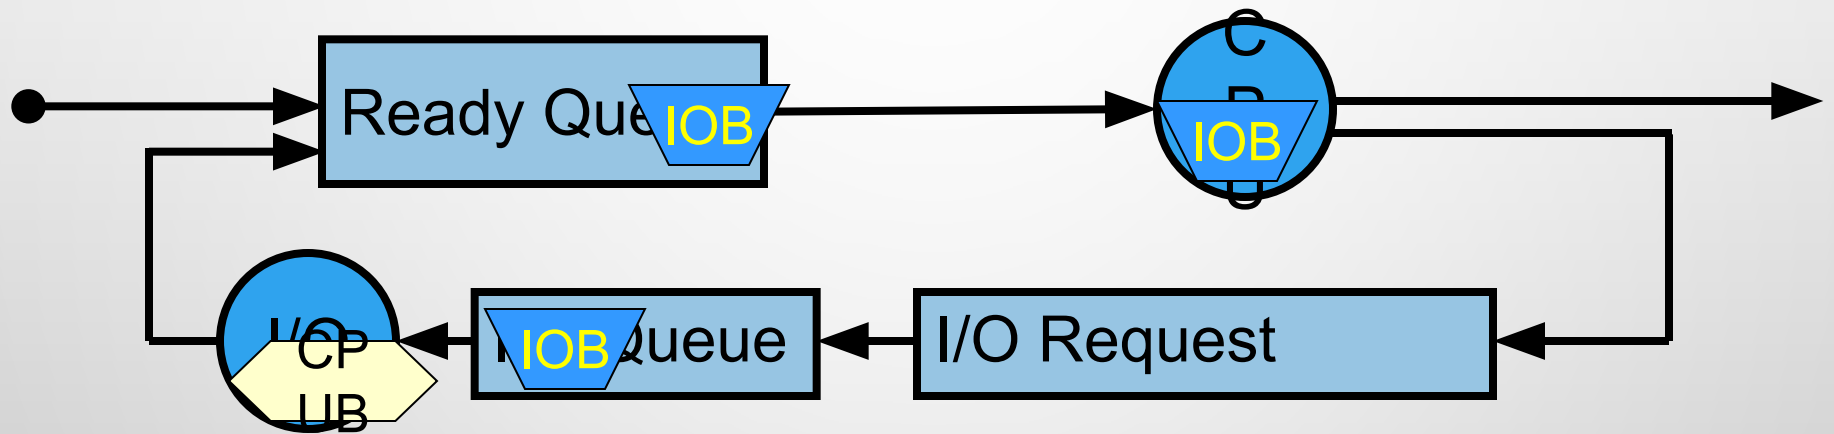
CPU is running CPUB

CONVOY EFFECT



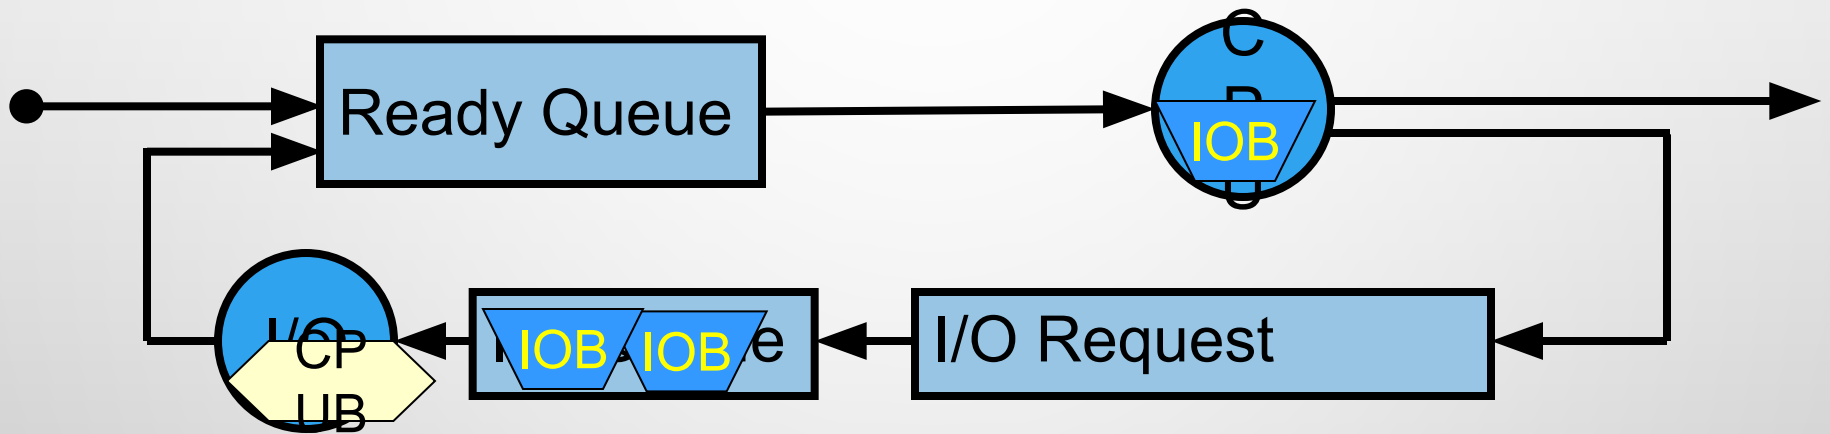
*CPU is running CPUB
I/O devices idle*

CONVOY EFFECT



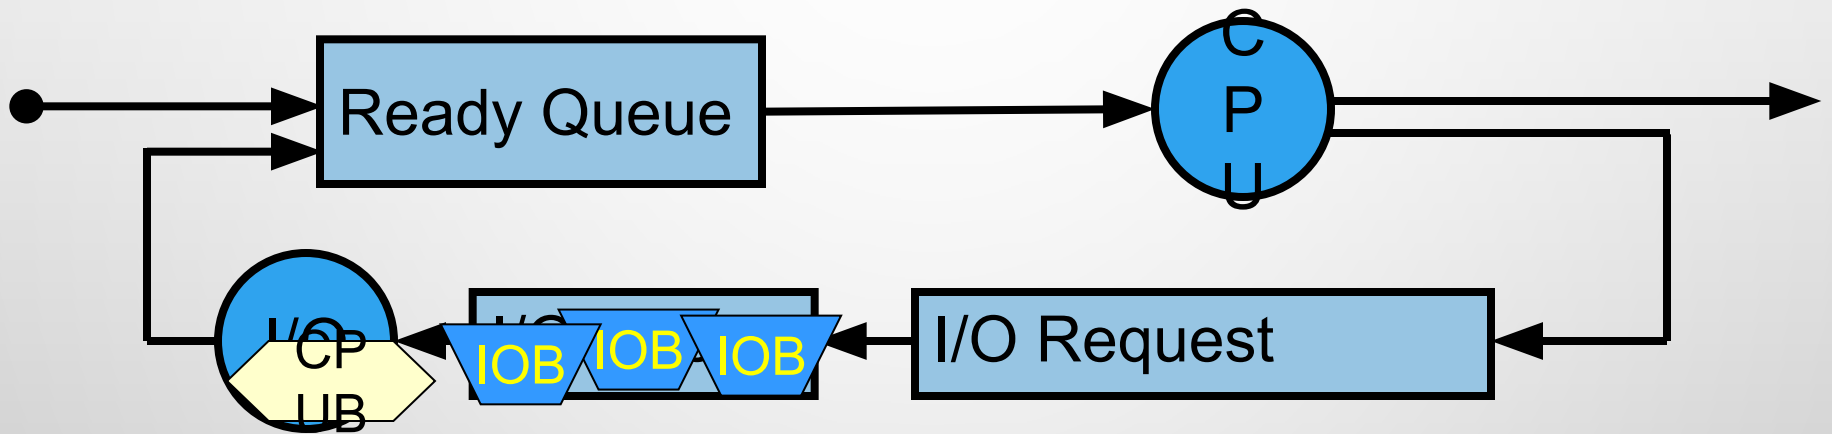
CPUB moves to I/O device

CONVOY EFFECT



I/O Bound jobs take very small amount of CPU time and go for I/O

CONVOY EFFECT




CPU idle

SHORTEST JOB FIRST (SJF)

- Scheduling algorithm in **batch** systems
- Schedule the job with the shortest run time first
- Requirement: **the run time needs to be known in advance**
- SJF is **optimal** in terms of turnaround, if **all** jobs arrive at **same time**

SJF: EXAMPLE

Process	Duration	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



Do it yourself

P4 turnaround: 3
P1 turnaround: 9
P3 turnaround: 16
P2 turnaround: 24

Total execution time: 24
The average turnaround:
 $(3+9+16+24)/4 = 13$

COMPARING TO FCFS

Process	Duration	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



P1 turnaround: 6
P2 turnaround: 14
P3 turnaround: 21
P4 turnaround: 24

The total time is the same.
The average turnaround:
 $(6+14+21+24)/4 = 16.25$
(comparing to 13)

SJF IS NOT ALWAYS OPTIMAL

- SJF OPTIMAL ONLY IF ALL JOBS HAVE ARRIVED AT SCHEDULING TIME

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 turnaround: 10

P2 turnaround: 10


The average turnaround (AWT):
 $(10+10)/2 = 10$

PREEMPTIVE SJF

- Also called **Shortest Remaining Time Next**
 - Schedule the job with the shortest **remaining** time required to complete
 - When **new** job **arrives**, compare its **total** time with the **remaining** time of the running job
 - If the new job needs less time the current job is suspended and the new job started
- Requirement: the run time needs to be known in advance

PREEMPTIVE SJF: SAME EXAMPLE

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 turnaround: 12
P2 turnaround: 2

The average turnaround:
 $(2+12)/2 = 7$

PROBLEM WITH PREEMPTIVE SJF?

- Starvation

- In some condition, a job is waiting for ever
- Example: Preemptive SJF
 - Process A with run time of 1 hour arrives at time 0
 - But every 1 minute, a short process with run time of 1 minute arrives
 - Result of Preemptive SJF: A never gets to run

INTERACTIVE SYSTEM

- Example: Servers
 - Serve multiple remote users all of whom are in a big hurry
- Performance Criteria
 - Min response time:
 - amount of time it takes from when a request was submitted until the **first response** is produced, not output
 - respond to requests quickly

INTERACTIVE SYSTEM

- Algorithms used here usually preemptive
 - Time is **sliced** into quantum (time intervals)
 - Scheduling decision is also made at the beginning of each quantum
- Representative algorithms:
 - Round-robin
 - Priority-based
 - Shortest process time
 - Guaranteed Scheduling
 - Lottery Scheduling
 - Fair Sharing Scheduling


ROUND ROBIN

- **Round Robin (RR)**

- Often used for timesharing
- Each process is given a time slice called a *quantum*
- It is run for the quantum or until it blocks
- RR allocates the CPU uniformly (fairly) across participants from ready queue.

- Problem:

- Do not consider priority
- Context switch overhead



IMPLEMENTING ROUND ROBIN

- Keep the ready queue as a FIFO queue of processes.
- **New** processes are added to the **tail** of the ready queue.
- The scheduler
 - picks the **first** process from the ready queue
 - sets a timer to interrupt after 1 time quantum, and
 - Starts the process.
- When the quantum is over
 - The running process will be put at the **tail** of the ready queue.

RR WITH TIME QUANTUM = 20

<u>Process</u>	<u>Run Time</u>
----------------	-----------------


P_1	53
-------	----

P_2	17
-------	----

P_3	68
-------	----

P_4	24
-------	----

- All processes arrive at time 0
- The **Gantt** chart is



- Higher average turnaround than SJF
- But better response time

RR: CHOICE OF TIME QUANTUM

- Performance depends on length of the timeslice
 - Context switching isn't a free operation.
 - If timeslice time is set too high
 - attempting to amortize context switch cost, you get FCFS.
 - i.e. processes will finish or block before their slice is up anyway
 - Poor response time
 - If it's set too low
 - you're spending all of your time context switching between threads.

PRIORITY SCHEDULING

- Each job is assigned a priority
- Select **highest** priority job to run next
- Rational: higher priority jobs are more important
 - Example: simulation vs. auto save a document
- Problems:
 - Low priority process may **starve**
- Solution:
 - Priority need to be **adjusted** depending on the situation

ASSIGN PRIORITY

- Two approaches
 - Static (for system with well known and regular application behaviors)
 - Dynamic (otherwise)
- Priority may be based on:
 - Cost to user.
 - Importance of user
 - Percentage of CPU time used in last X hours

EXAMPLE: DYNAMIC PRIORITY ASSIGNMENT

- Whenever highly I/O bound processes wants the CPU it should be given the CPU immediately.
- Why?
- A simple algorithm for giving priority to I/O bound processes is to set the priority to $1/f$
 - f is the fraction of the last quantum used by a process
 - A process that used only 1 msec of its 50 msec quantum would get priority 50
 - A process that used 25 msec of its 50 msec quantum would get priority 2

SHORTEST PROCESS NEXT

- Let's apply SJF for **interactive** processes
 - General pattern of a **interactive** process: CPU burst, I/O burst, ...
 - Let's regard the execution of each CPU burst as a separate "**job**"
 - Now we can minimize overall response time by running the process with shortest "**job**" first

SHORTEST PROCESS NEXT

- How to know the length of the next CPU burst?
 - A possible answer: Exponential averaging
 - Make **estimate** based on past behavior and run the process with the shortest estimated CPU burst
 - Let the **current estimated** CPU burst is τ_n
 - length of the nth CPU burst t_n
 - predicted value for the next CPU burst $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \cdots + (1 - \alpha)^j \alpha t_{n-j} + \cdots + (1 - \alpha)^{n+1} \tau_0.$$

GUARANTEED SCHEDULING

- Make promises to users about performance & then meet those promises
- With n processes running, each one should get $1/n$ of the CPU cycles
- Calculate **ratio** for each process
 - $$\frac{\text{Amount of CPU time process has had since its creation}}{\text{Amount of CPU time process **should have** since creation}}$$
- Run the process with the **lowest** ratio until its ratio has moved above its closest competitor
- Problem:
 - Implementation is **difficult**

LOTTERY SCHEDULING

- Give processes lottery **tickets** for CPU time
- Whenever a scheduling decision has to be made, a lottery ticket is chosen **at random**, and the process holding that ticket gets the CPU
- More important processes can be given extra tickets, to increase their chances of winning.
- If there are 100 tickets and one process holds 20 of them, it will have a 20% chance of winning each lottery
- In the long run, it will get about 20% of the CPU
- Highly Responsive:
 - if a **new** process shows up and is granted some tickets
 - at the very next lottery it will have a chance of winning in proportion to the number of tickets it holds

LOTTERY SCHEDULING

- Cooperating processes may exchange tickets if they wish.
 - When a **client process** sends a message to a **server process** and then blocks, it may give **all** of its tickets to the server, to **increase** the chance of the **server** running next
 - After finishing, it returns the tickets so that the client can run again.
- Can solve problems that are **difficult** to handle with other methods
 - In a video server several processes are feeding video **streams** to their clients, but at **different** frame rates.
 - Let the processes need frames at 10, 20, and 25 frames/sec.
 - By allocating these processes 10, 20, and 25 tickets, respectively, they will automatically divide the CPU in approximately the correct proportion, that is, 10:20:25.

REAL-TIME SYSTEMS

- **Time** plays an essential role
- Usually the computer must react appropriately to **events** generated by external devices within a **fixed** amount of time
 - patient monitoring in a hospital intensive-care unit,
 - the autopilot in an aircraft
 - robot control in an automated factory
- Getting right answer but too late == Getting **nothing** at all
 - may have **catastrophic** consequences
 - financial loss
 - major equipment damage
 - loss of life

The image features a light gray background with a subtle gradient. In the top-left and bottom-right corners, there are several realistic water droplets of varying sizes. These droplets are rendered with soft shadows and highlights, giving them a three-dimensional appearance. The text "THANKS FOR YOUR PATIENCE" is centered in the middle of the image in a bold, black, sans-serif font.

THANKS FOR YOUR PATIENCE